

# Less Spooky Containers with Top Security Best Practices



**ALVARO IRADIER**

Product Analyst, Sysdig



# Agenda

- Container Security Model
- Why: shifting left security
- Best practices
  - Build
  - Distribute
  - Deploy
  - Run



# Container security model

# Containers!

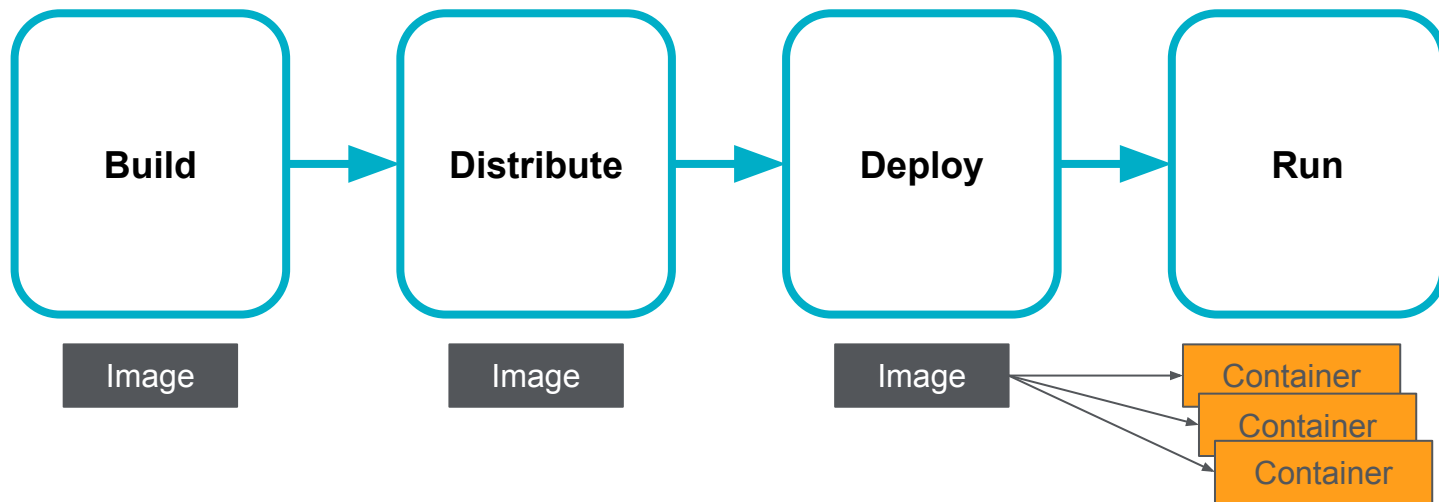




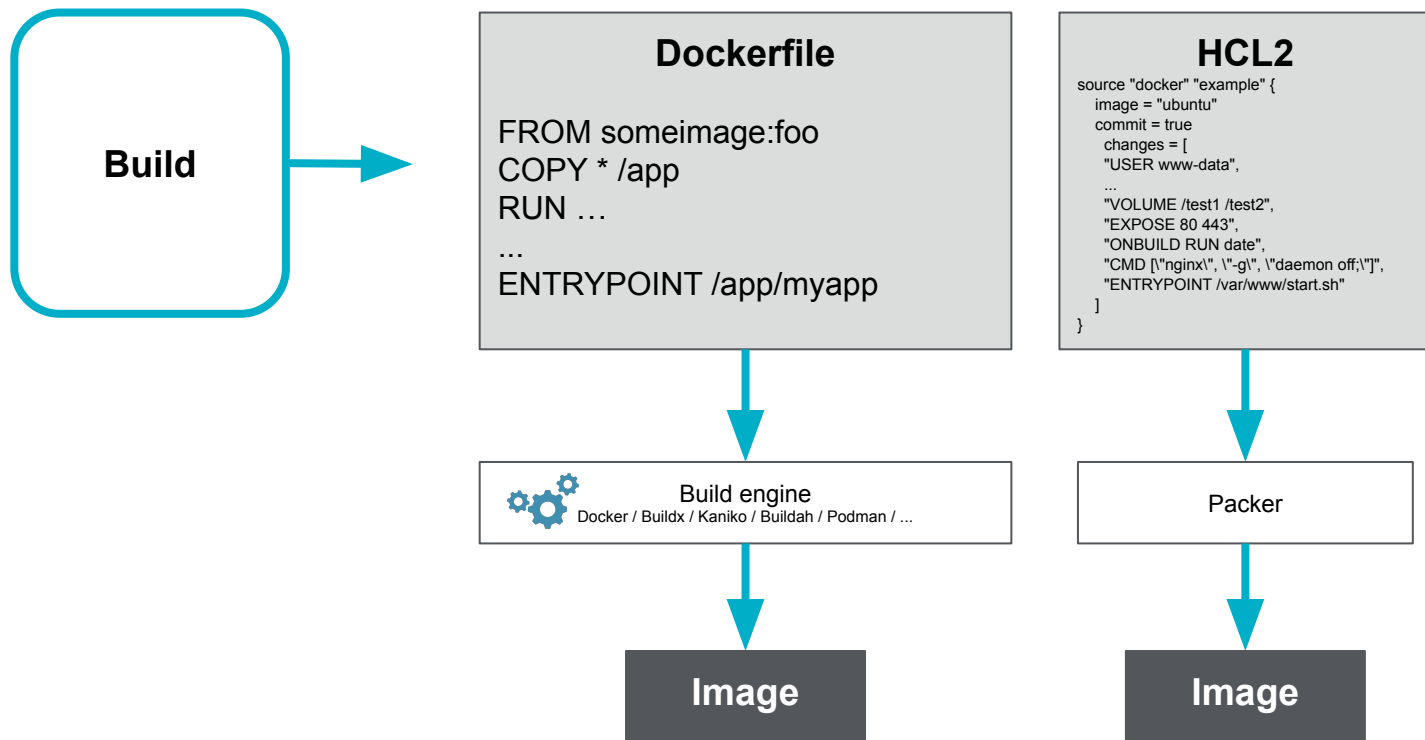
# Also containers!





# Container life stages



# Building an image




# Distributing an image






# Deploying a container



# Running a container



## Container

- Image root filesystem + chroot
  - Contains application(s) + dependencies
- Configuration:
  - Entrypoint and arguments
  - Environment variables
  - ...
- Isolation: namespaces (PID, IPC, MOUNTS, UTS, NET, ...)
- Resource limits: CGroups
- Privilege limits: Seccomp, AppArmor, ...



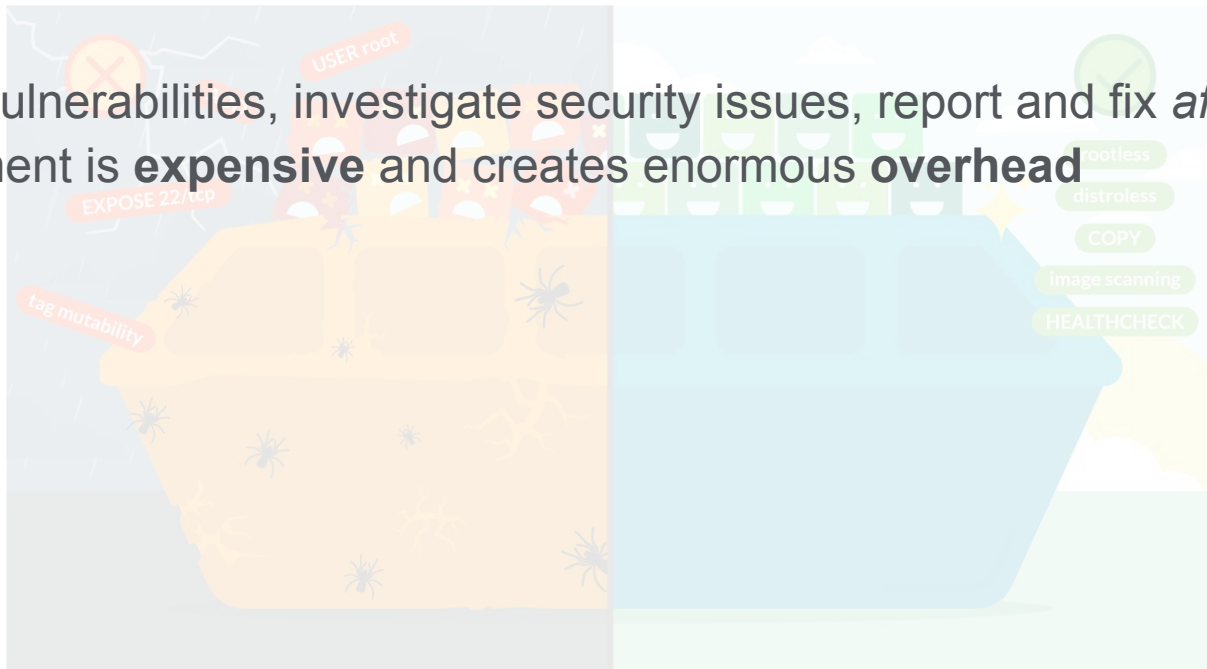
# Why: shifting left security

# Why



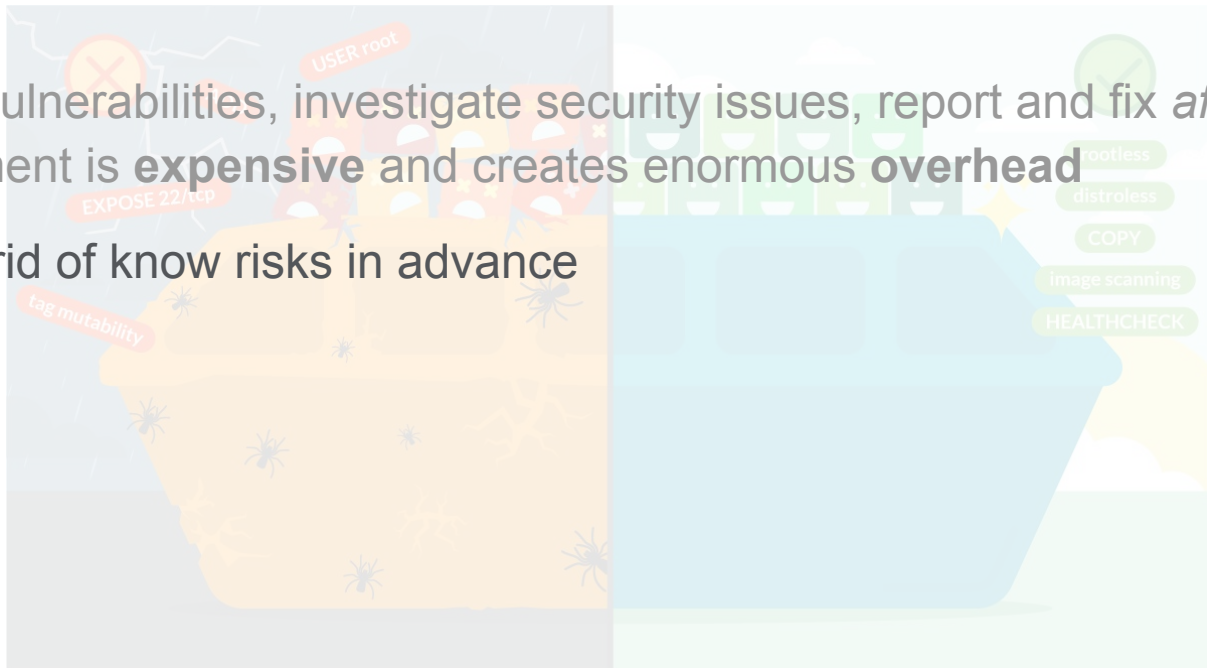
# Why: shifting left security

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**



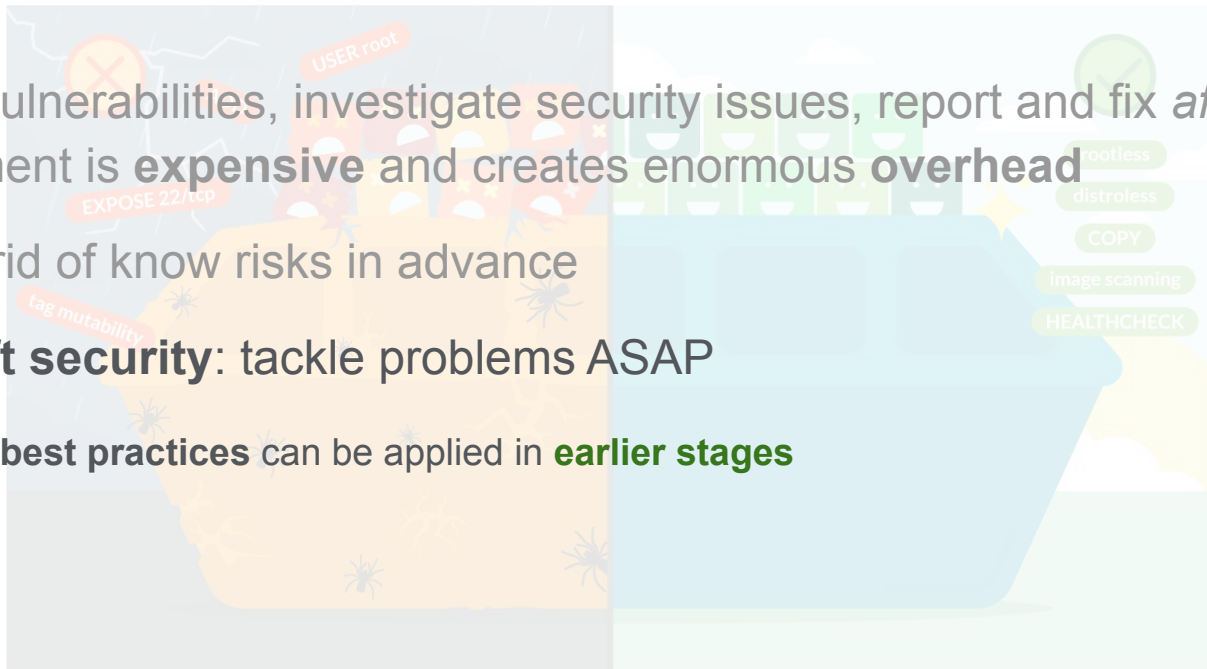
# Why: shifting left security

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of know risks in advance



# Why: shifting left security

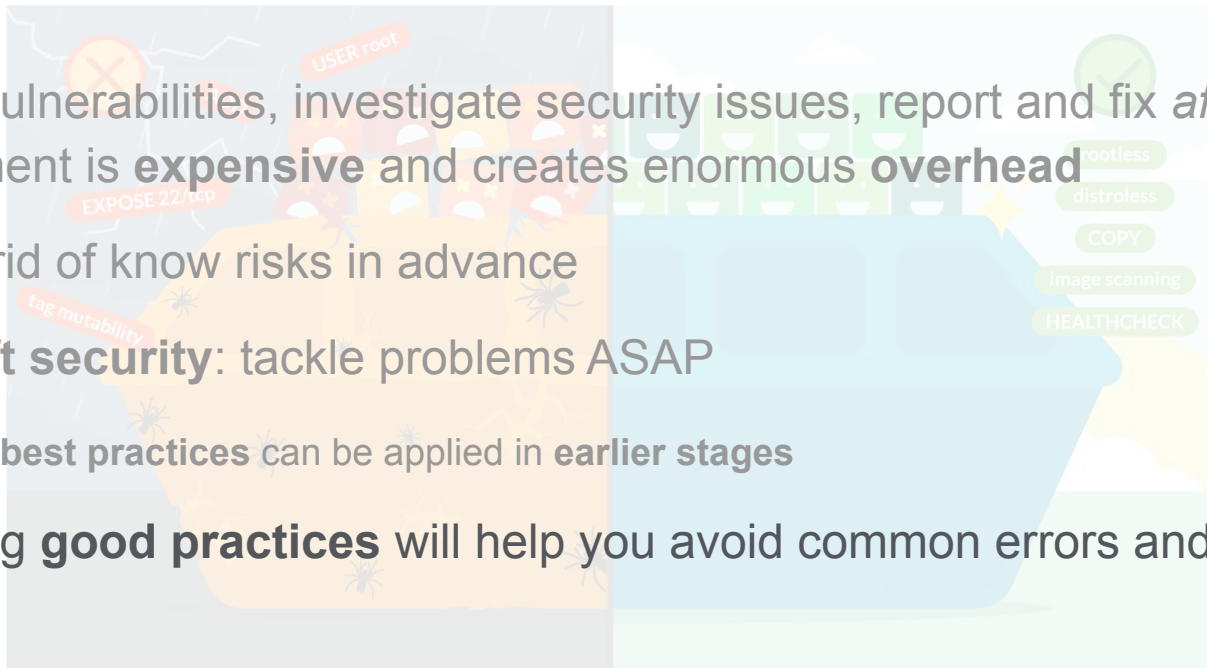
- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of known risks in advance
- **Shift left security**: tackle problems ASAP
  - **Build best practices** can be applied in **earlier stages**






# Why: shifting left security

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of known risks in advance
- **Shift left security**: tackle problems ASAP
  - **Build best practices** can be applied in **earlier stages**
- Following **good practices** will help you avoid common errors and pitfalls



# Shifting left security






# Best practices - build

# Unnecessary privileges

- **Problem:** containers often *running* with more privileges than required
  - Our [recent report](#) highlighted that 58% of images are running the container entrypoint as root (UID 0)



- **Risks:**
  - Attackers can exploit vulnerabilities or bugs to gain access to other services or resources
  - <https://sysdig.com/blog/lateral-movement-cloud-containers/>
  - Openshift or others blocking root containers



# Unnecessary privileges

- **Prevention:**
  - Follow the principle of **least privilege** so your service or application only has access to the resources and information necessary to perform its purpose
  - **USER directive**, run as non-root by default

# Unnecessary privileges

- Step 1: Initial Dockerfile

```
FROM alpine
COPY ./src/example_app /example_app
ENTRYPOINT /example_app
```

```
> docker build -f Dockerfile1 -t test1_1 .
[+] Building 0.5s (7/7) FINISHED

> docker run --name test --rm -p 5000:5000 test1_1
Listening at :5000
```

# Unnecessary privileges

- Testing the app

```
> curl "localhost:5000/?name=webinar"
Hi webinar!
> curl "localhost:5000/?name=alvaro"
Hi alvaro!
> docker exec -ti test sh
/ # ls -lh
-rw-r--r--      1 root      root           111 Apr  7 11:36 access.log
-rwxr-xr-x      1 root      root        5.9M Apr  7 11:27 example_app
...
/ # cat access.log
[2021-04-07T11:36:14Z] IP=172.17.0.1:60974 name=webinar
[2021-04-07T11:36:16Z] IP=172.17.0.1:60978 name=alvaro
/ # ps
PID   USER      TIME  COMMAND
    1  root      0:00  /example_app
```



# Unnecessary privileges

- Step 2: running as non-root

```
FROM alpine
COPY ./src/example_app /app/example_app
WORKDIR /app
USER 1000
ENTRYPOINT /app/example_app
```

```
> docker build -f Dockerfile2 -t test1_2 .
[+] Building 0.5s (7/7) FINISHED
> docker run --name test --rm -p 5000:5000 test1_2
panic: open access.log: permission denied
```

# Unnecessary privileges

- Step 3: fix permissions

```
FROM alpine
COPY ./src/example_app /app/example_app
WORKDIR /app
RUN chown 1000:1000 /app
USER 1000
ENTRYPOINT /app/example_app
```

```
> docker build -f Dockerfile3 -t test1_3 .
[+] Building 0.5s (7/7) FINISHED
> docker run --name test --rm -p 5000:5000 test1_3
Listening at :5000
```

# Unnecessary privileges

- Testing the **rootless** app

```
...  
> docker exec -ti test sh  
/app $ cat access.log  
[2021-04-07T11:45:56Z] IP=172.17.0.1:60996 name=webinar  
[2021-04-07T11:45:58Z] IP=172.17.0.1:61000 name=alvaro  
/app $ ls -lh  
total 6M  
-rw-r--r--      1 1000      root           111 Apr  7 11:45 access.log  
-rwxr-xr-x      1 root      root           5.9M Apr  7 11:27 example_app  
/app $ ps  
PID   USER      TIME  COMMAND  
  1  1000      0:00  /app/example_app
```

# Attack surface

- **Problem:** including unnecessary packages or exposing unused ports
- **Risks:**
  - Your system is more exposed to attacks
  - Using components not under your control
- **Prevention:**
  - Reduce attack surface
  - Protect and keep under control

# Attack surface: reduce attack surface

- **Don't:** Use big, generic distro images if not needed (i.e. ubuntu)

```
> docker run ... quay.io/sysdig/secure-inline-scan:2 image-ubuntu -k $SYSDIG_SECURE_TOKEN --storage-type docker-daemon
Inspecting image from Docker daemon -- distroless-1:latest
Full image: docker.io/library/image-ubuntu
Full tag: docker.io/library/image-ubuntu:latest
...
Analyzing image..
Analysis complete!
...
Evaluation results
- warn dockerfile:instruction Dockerfile directive 'HEALTHCHECK' not found, matching condition 'not_exists' check
- warn dockerfile:instruction Dockerfile directive 'USER' not found, matching condition 'not_exists' check
- warn files:suid_or_guid_set SUID or SGID found set on file /bin/mount. Mode: 0o104755
- warn files:suid_or_guid_set SUID or SGID found set on file /bin/su. Mode: 0o104755
- warn files:suid_or_guid_set SUID or SGID found set on file /usr/bin/chage. Mode: 0o102755
...
Vulnerabilities report
  Vulnerability  Severity  Package                                Type  Fix version  URL
- CVE-2019-18276  Low      bash-4.3-14ubuntu1.4                  dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2019-18276
- CVE-2016-2781   Low      coreutils-8.25-2ubuntu3~16.04         dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781
- CVE-2017-8283   Negligible dpkg-1.18.4ubuntu1.6                  dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2017-8283
- CVE-2020-13844  Medium   gcc-5-base-5.4.0-6ubuntu1~16.04.12    dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2020-13844
...
- CVE-2018-20839  Medium   systemd-sysv-229-4ubuntu21.29         dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-20839
- CVE-2016-5011   Low      util-linux-2.27.1-6ubuntu3.10         dpkg  None         http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-5011
...
```

## Attack surface: reduce attack surface

- **Don't:** Use big, generic distro images if not needed (i.e. ubuntu)

- Things most likely won't need in your final image:

- gcc-5 compiler
- sysV compatibility
- dpkg? bash?

- More than 100 vulnerabilities detected!

- **Don't:** Official images might not be the best fit *per se*

- Regarding security and minimalism, they might not be updated that often and can include extra packages for general use cases

```
> docker build -f Dockerfile --no-cache --rm -t image-ubuntu -k $$SYSDIG_SECURE_TOKEN --storage-type docker-daemon
Inspecting image...
Full image: docker.io/library/image-ubuntu
Full tag: localhost:5000/image-ubuntu
Analyzing image...
Analysis completed
...
Evaluation results
- warn dockerfile:instruction Dockerfile directive 'HEALTHCHECK' not found, matching condition 'not_exists' check
- warn dockerfile:instruction Dockerfile directive 'USER' not found, matching condition 'not_exists' check
- warn file:suid_or_guid_set SUID or SGID found set on file /bin/mount. Mode: 0o14755
- warn file:suid_or_guid_set SUID or SGID found set on file /bin/mount. Mode: 0o14755
- warn files:suid_or_guid_set SUID or SGID found set on file /usr/bin/chage. Mode: 0o102755
...
Vulnerabilities report
Vulnerability Severity Package Type Fix version URL
- CVE-2019-18276 Low bash-4.3-14ubuntu1.4 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2019-18276
- CVE-2016-2781 Low dpkg-1.18.2ubuntu1 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781
- CVE-2017-8283 Low dpkg-1.18.2ubuntu1 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2017-8283
- CVE-2020-13844 Medium gcc-5-base-5.4.0-6ubuntu1~16.04.12 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2020-13844
...
- CVE-2018-20839 Medium systemd-sysv-229-4ubuntu2.19 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-20839
- CVE-2016-5011 Low mail-2.7.3-6ubuntu1.19 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-5011
...
```

# Attack surface: reduce attack surface

- **Be minimal** in your base base
  - *alpine* versions
  - FROM scratch
  - Distroless (<https://github.com/GoogleContainerTools/distroless>):
    - i.e.: FROM `gcr.io/distroless/base-debian10`
      - Basic set of packages, including just required libraries like *glibc*, *libssl*, and *openssl*.
    - Slimmer: FROM `gcr.io/distroless/static-debian10`
      - For statically compiled applications like Go that don't require libc



# Attack surface: reduce attack surface

- Check for **optimized** vs generic versions
  - **Example:** [bitnami/node](#) vs [official node image](#)
    - customized versions on top of a minideb distribution
    - frequently updated with the latest bug fixes
    - signed with Docker Content Trust
    - pass a [security scan for tracking known vulnerabilities](#)

# Attack surface: build, size & reproducibility

- **Don't:**
  - Build the application externally, copy into the container
  - Bad reproducibility

```
FROM alpine
COPY ./src/example_app /example_app
USER 1000
WORKDIR /tmp
ENTRYPOINT /example_app
```

# Attack surface: build, size & reproducibility

- **Don't:**
  - **Build directly inside the final container**
    - Big image size and multiple layers, build toolchain included in the container, other unrequired packages, remainings of the application source code, ...

```
FROM alpine
RUN apk add go
COPY ./src/ /src
WORKDIR /src
RUN go build .
ENTRYPOINT /src/example_app
```

# Attack surface: build, size & reproducibility

- **Multistage builds:**

```
#This is the "builder" stage
```

```
FROM golang:1.16 as builder
```

```
WORKDIR /my-go-app
```

```
COPY src .
```

```
RUN GOOS=linux GOARCH=amd64 go build .
```

```
#This is the final stage, and we copy artifacts from "builder"
```

```
FROM alpine
```

```
COPY --from=builder /my-go-app/example_app /bin/example_app
```

```
ENTRYPOINT ["/bin/example_app"]
```

# Attack surface: build, size & reproducibility

- **Multistage builds:**
  - Reproducible builds, always same build environment
  - Minimal image size, no build tools or undesired packages

```
> docker images
test2_1          latest          681b9e590ae9   5 minutes ago   448MB
test2_2          latest          37d02efde1e4   35 seconds ago  7.55MB
...
```

# Attack surface: build, size & reproducibility

- **Example: multistage build with nodejs and typescript**

(<https://github.com/kevinpollet/typescript-docker-multi-stage-build>):

```
FROM node:14-alpine AS builder
WORKDIR /usr/src/app
COPY typescript-docker-multi-stage-build/package*.json ./
RUN npm ci
COPY typescript-docker-multi-stage-build/tsconfig*.json ./
COPY typescript-docker-multi-stage-build/src src
RUN npm run build
```

```
FROM node:14-alpine
ENV NODE_ENV=production
WORKDIR /usr/src/app
RUN chown node:node .
USER node
COPY typescript-docker-multi-stage-build/package*.json ./
RUN npm install
COPY --from=builder /usr/src/app/lib/ lib/
EXPOSE 3000
ENTRYPOINT [ "node", "lib/server.js" ]
```

# Attack surface: trusted sources

- **Don't:** Use `docker.io/johndoe hacker/mycustom-node-image:latest`

```
FROM docker.io/johndoe hacker/mycustom-node-image:latest
...
```

- Inherit all of the problems and vulnerabilities from that image
- Who builds and publishes that image?
- Is it updated regularly?
- How is it built?
- Are we sure the published version is really from the public Dockerfile?



# Attack surface: trusted sources

- Prefer **verified** and **official** images from **trusted repositories and providers**
  - A common choice: [Redhat UBI](#) (Universal Base Images)
  - Enterprise support, upgrades or certifications vs minimalism
- When using custom images, check for the **image source** and the Dockerfile, and **build your own base image**

# Credentials & Confidentiality

- **Problem:** leaking credentials or confidential information in your images
- **Risks:**
  - Attackers can use leaked credentials to access your systems
  - Exposure of confidential or sensitive information
- **Prevention:**
  - Don't include sensitive data in the image, use external stores
  - Good practices to avoid data leaks

# Credentials & Confidentiality: sensitive data

- **Don't:**
  - Include hard coded credentials
  - Add credentials file or environment variables

```
FROM alpine
...
ENV SECURE_API_TOKEN=ajhda8-12312-29889234-foo
COPY aws_credentials /home/app/.aws/credentials
...
ENTRYPOINT /example_app
```

# Credentials & Confidentiality: sensitive data

```
> docker inspect test3_1
[
  {
    "Id": "sha256:18440a2433ea49efa686febc6f02c21a652a498523ed42e00cf79ebf3717cc0a",
    "RepoTags": [
      "test3_1:latest"
    ],
    "RepoDigests": [],
    "Parent": "",
    "Comment": "buildkit.dockerfile.v0",
    "Created": "2021-04-07T17:01:00.8254965Z",
    "Container": "",
    "ContainerConfig": {
      ...
    },
    "DockerVersion": "",
    "Author": "",
    "Config": {
      ...
      "Env": [
        "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin",
        "SECURE_API_TOKEN=ajhda8-12312-29889234-foo"
      ]
    }
  }
]
```

# Credentials & Confidentiality: sensitive data

```
> docker run --entrypoint /bin/sh --rm test3_1 -c "cat /home/app/.aws/credentials"
[default]
aws_access_key_id = SOME-ACCESS-KEY
aws_secret_access_key = SOME-SECRET-KEY
```

# Credentials & Confidentiality: layers

- **Even if the file is removed!! - Don't** forget the **layered** nature of images (each command creates a new layer)
  - Removing a file in a layer still takes space and file can be accessed

```
FROM alpine
...
COPY aws_credentials /home/app/.aws/credentials
...
RUN rm /home/app/.aws/credentials
ENTRYPOINT /example_app
```

# Credentials & Confidentiality: layers

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory
>
```

# Credentials & Confidentiality: layers

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

>
```



# Credentials & Confidentiality: layers

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

> cat test3_2/index.json
{"schemaVersion":2,"manifests":[{"mediaType":"application/vnd.oci.image.manifest.v1+json","digest":"sha256:a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089","size":657}]}
```

# Credentials & Confidentiality: layers

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

> cat test3_2/index.json
{"schemaVersion":2,"manifests":[{"mediaType":"application/vnd.oci.image.manifest.v1+json","digest":"sha256:a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089","size":657}]}

> cat test3_2/blobs/sha256/a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089 | jq
{
  "schemaVersion": 2,
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "digest": "sha256:6a1057f9fe2693956a5bb40bd9e3ee624171f4145dc5b6c7d83b03e4d2774688",
    "size": 1236
  },
  "layers": [
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:2d5a20755f17e53a78fdfeebffff1100a88ec7941c727a9538932b0409ca7bf5c",
      "size": 2899855
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7"
```

# Credentials & Confidentiality: layers

```
> cat test3_2/blobs/sha256/a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089 | jq
{
  "schemaVersion": 2,
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "digest": "sha256:6a1057f9fe2693956a5bb40bd9e3ee624171f4145dc5b6c7d83b03e4d2774688",
    "size": 1236
  },
  "layers": [
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:2d5a20755f17e53a78fdfeebfff1100a88ec7941c727a9538932b0409ca7bf5c",
      "size": 2899855
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7",
      "size": 284
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:27bc864279bbb3afc3ef38c2acd34b6c2897249e55076e518e7abc36f9ec4ec3",
      "size": 199
    }
  ]
}
```

# Credentials & Confidentiality: layers

```
> tar xvfz test3_2/blobs/sha256/f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7
x home/
x home/app/
x home/app/.wh..wh..opq
x home/app/.aws/
x home/app/.aws/credentials

>
```

# Credentials & Confidentiality: layers

```
> tar xvfz test3_2/blobs/sha256/f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7
x home/
x home/app/
x home/app/.wh..wh..opq
x home/app/.aws/
x home/app/.aws/credentials

> cat home/app/.aws/credentials
[default]
aws_access_key_id = SOME-ACCESS-KEY
aws_secret_access_key = SOME-SECRET-KEY
```

# Credentials & Confidentiality: leaks

- **Don't:** Leak files from the build context

```
docker build -t myimage .
```

- The “.” parameter is the build context
- All the files in the build context are sent to the docker daemon
  - You can copy confidential or unnecessary files into the container, like configuration files, credentials, backups, lock files, temporary files, sources, subfolders, dotfiles, etc.
- COPY and ADD commands work from the build context

# Credentials & Confidentiality: leaks

- **Example**

```
docker build -f Dockerfile -t myimage .
```

```
...  
COPY . /my-app
```

This would copy everything inside the build context, which for the “.” example, includes the Dockerfile itself.

# Credentials & Confidentiality: leaks

Good practices:

- Use a **clean build context**

```
docker build -f Dockerfile -t myimage files/
```

- Use **.dockerignore** file
- Prefer **COPY over ADD**, and **avoid wildcards**
  - COPY is more explicit, more predictable and less error prone
  - ADD can add files from a URL or from a .tar file



# Linting

- Tools like [Haskell Dockerfile Linter](#) (hadolint) can detect bad practices in your Dockerfile, and even expose issues inside the shell commands executed by the RUN instruction.
- Image scanners (like Sysdig's) are also capable of detecting bad practices via customizable rules
- **Automate**: Consider incorporating such a tool in your CI pipelines.

IMAGE SCANNING  
Policies > Edit Policy

Default policies are read-only. Duplicate this policy to make changes.

Name: Default Configuration Policy - Dockerfile Best Practices

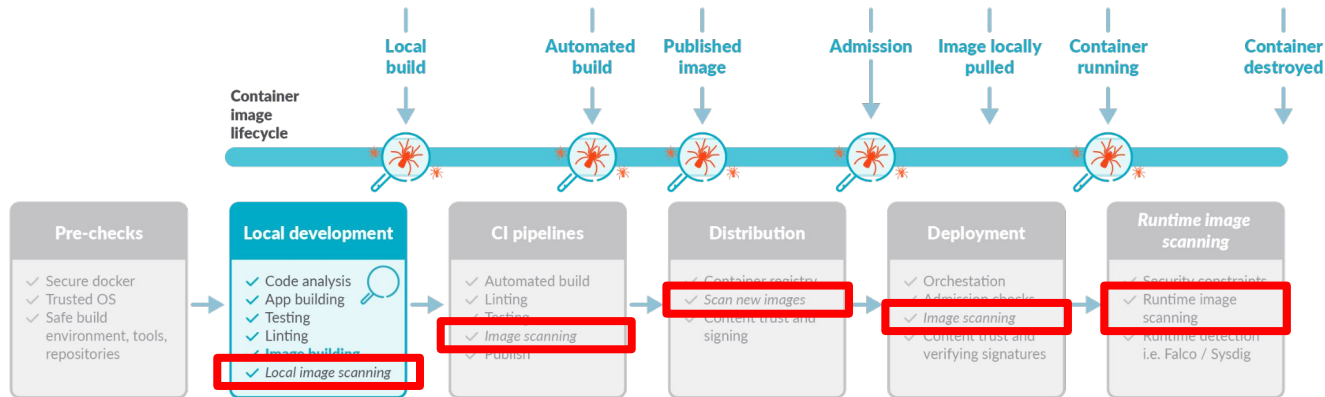
Description: This policy provides out of the box rules around Dockerfile best practices. We frequently update these policies and if you'd like to modify the policy you should use this as a base template to avoid modifications being overwritten.

Rules				
Vulnerabilities	Stale feed data	Max days since sync: 7	Warn	X
Dockerfile	Instruction	Instruction: RUN; Check like: Value: *apt-get upgrade.*	Warn	X
Dockerfile	Instruction	Instruction: RUN; Check like: Value: *yum upgrade.*	Warn	X
Dockerfile	Instruction	Instruction: HEALTHCHECK; Check: not_exists	Warn	X
Dockerfile	Effective user	Type: blacklist; Users: root	Warn	X
Dockerfile	Exposed ports	Type: blacklist; Ports: 22	Warn	X
Dockerfile	Instruction	Instruction: LABEL; Check: !=; Value: latest	Warn	X
Dockerfile	Instruction	Instruction: ENV; Check like: Value: *passwordPASSWORDpasswordPASSWORDSecretSECRET*	Warn	X
Dockerfile	Instruction	Instruction: USER; Check: not_exists	Warn	X
Dockerfile	Instruction	Instruction: ADD; Check: exists	Warn	X

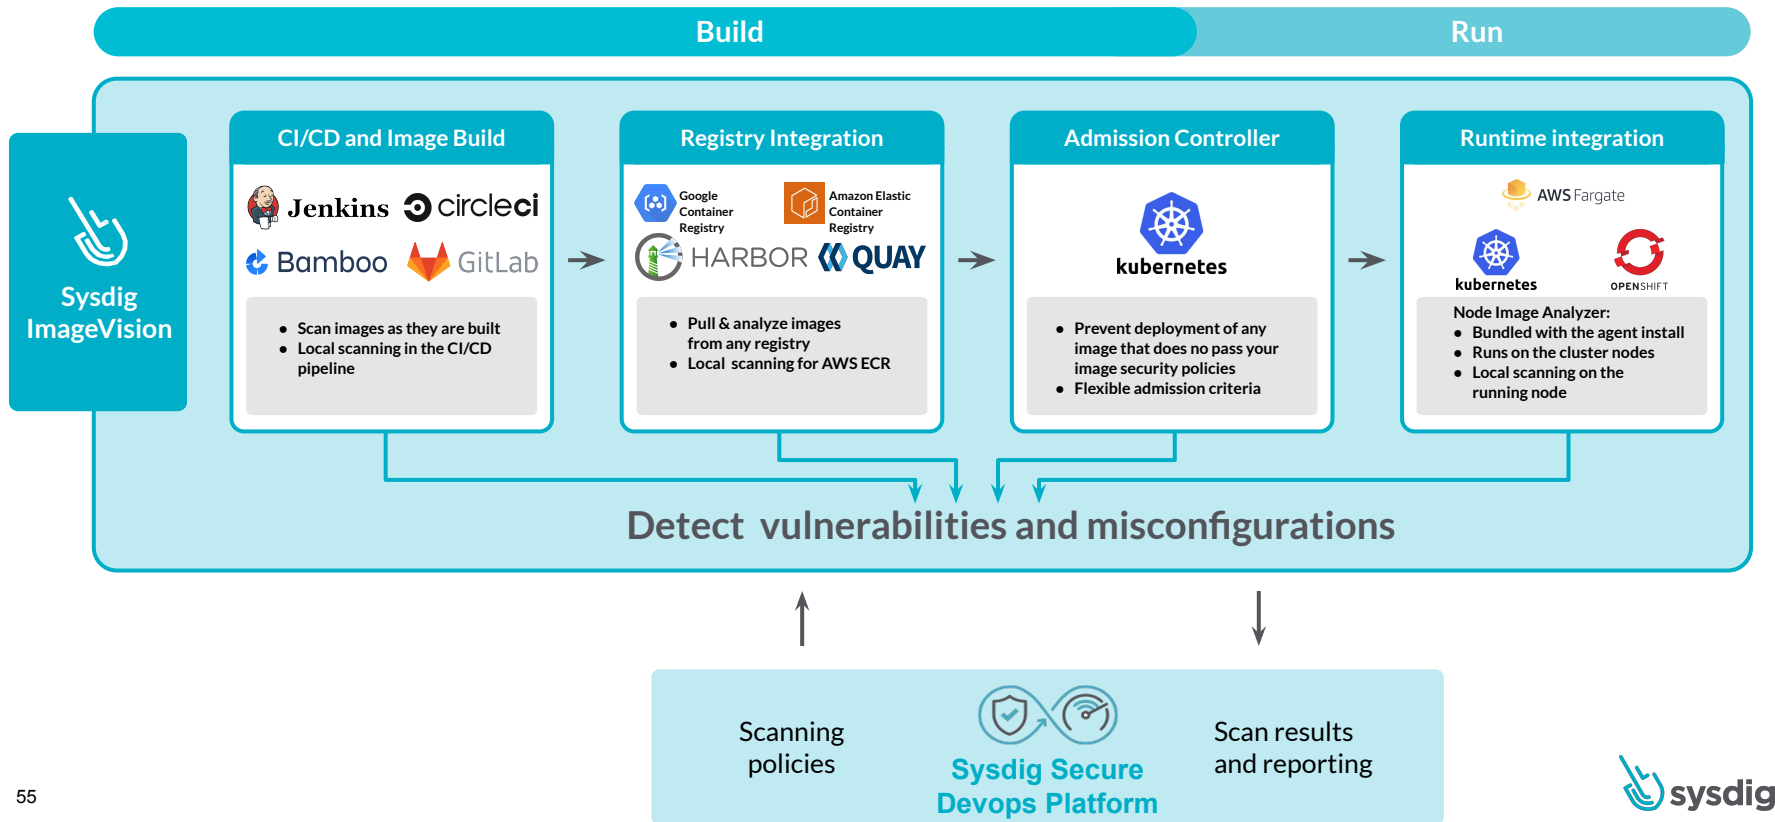
Dockerfile	Effective user	Type: blacklist; Users: root
Dockerfile	Exposed ports	Type: blacklist; Ports: 22

# Image Scanning

- **Image scanning** can be implemented at different stages
  - Detect bad practices and known vulnerabilities
  - The earlier the scan is performed, the better



# Image Scanning





# Best practices - distribute

# Image distribution

- **Problems:**

- Images pushed to the registry from untrusted sources
- Images bypassing security scans or best practices at build
- Images replaced or tampered

- **Risks:**

- Unsecure images deployed and running in your environment

# Image distribution: protect registry

- Protect registry access
  - Proper credentials and permissions
  - Internal registry (don't expose to the Internet)
- Protect communications
  - HTTPS / TLS
  - Valid certificates

# Image distribution: enforce good practices

- Image scanning when a new image is **pushed** to the registry
  - Report vulnerabilities and bad practices
  - Integrated image scanning (i.e. Harbor, ECR, ...)
  - Connect with external tools (Sysdig)
- Harbor and others can block vulnerable images - no pull allowed
- Full registry scan: security posture

# Image distribution: protect content

- Use tools to ensure images are not tampered via digital signatures:
  - Docker content trust + Docker notary or Harbor notary
  - Sigstore + cosign - <https://www.sigstore.dev/>
- Verify image signatures at deployment time
  - i.e. Connaisseur admission controller in Kubernetes
  - Trust configuration for podman, cri-o, etc:  
**/etc/containers/policy.json**





# Best practices - deploy

# Unnecessary privileges: UID

- Remember **unnecessary privileges**?
  - The **orchestrator** or **runtime** environment (i.e., docker run, kubernetes, etc.) has the last word on who is the running container effective user.
- **Prevention**
  - Avoid **running your containers as root**. Allow running as **any UID**
    - Openshift and some Kubernetes clusters will apply **restrictive policies by default**, **preventing root** containers from running or using a **random UID**
    - Simplifies permissions with host mounts: match container and host UIDs

# Unnecessary privileges: UID

- Try to un as UID "1001" instead of the default "1000"

```
> docker run -u 1001 --name test --rm -p 5000:5000 test1_3  
panic: open access.log: permission denied
```

```
FROM alpine  
COPY ./src/example_app /app/example_app  
USER 1000  
WORKDIR /tmp  
ENTRYPOINT /app/example_app
```

```
> docker run -u 1001 --name test --rm -p 5000:5000 test1_4  
Listening at :5000
```

# Unnecessary privileges: UID

- Verifying the container is running with UID 1001

```
...  
> docker exec -ti test sh  
/tmp $ ls -lh  
total 4K  
-rw-r--r--    1 1001      root          111 Apr  7 11:53 access.log  
/tmp $ ps  
PID   USER      TIME  COMMAND  
    1  1001      0:00  /app/example_app  
/tmp $ ls -lh /app  
total 5.9M  
-rwxr-xr-x    1 root      root      5.9M Apr  7 11:27 example_app
```

# Unnecessary privileges: perms. & capab.

- **Restrict application permissions and capabilities** on runtime
  - In case your container is compromised, the range of action available to an attacker is limited
  - [--cap-drop flag](#) in Docker
  - [securityContext.capabilities.drop](#) in Kubernetes
  - AppArmor in [Docker](#) or [Kubernetes](#)
  - Seccomp in [Docker](#) or [Kubernetes](#).

# Unnecessary privileges: kubernetes

```
apiVersion: v1
kind: Pod
...
metadata:
  ...
  annotations:
    container.apparmor.security.beta.kubernetes.io/hello: localhost/k8s-apparmor-example-deny-write
spec:
  securityContext:
    runAsUser: 1001
    runAsGroup: 1001
    fsGroup: 2000
    ...
  containers:
    - name: foo
      image: my-nice-app:2.0.3
      ...
      securityContext:
        runAsUser: 1002
        capabilities:
          drop:
            - all
          add: ["SYS_TIME"]
        seccompProfile:
          type: Localhost
          localhostProfile: my-profiles/profile-allow.json
```

# Unnecessary privileges: IaC

Loop from Production to Source

Detect runtime drift and instantly map it back to the IaC



Infra as code  
source file

```
apiVersion: v1
kind: Pod
metadata:
  name: java-app
spec:
  securityContext:
    runAsNonRoot: True
```



Prod cluster-tampered



Namespace.  
name = java-app

```
apiVersion: v1
kind: Pod
metadata:
  name: java-app
spec:
  securityContext:
    runAsNonRoot: False
```



<https://apolicy.io/>

# Unnecessary privileges: rootless containers


<https://rootlesscontaine.rs/>

- **Container runtime** (i.e. Docker daemon) with non-root user
  - Run inside a new user namespace
    - UID 0 *inside* the container is standard user *outside* the container
    - UID / GID mapping (via /etc/subuid and /etc/subgid)
  - Docker - <https://docs.docker.com/engine/security/rootless/>
  - Podman - <https://github.com/containers/podman#rootless>
  - Kubernetes - <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-in-users/>
  - Some shortcomings - due to limited privileges



# Mutant tags

- **Beware!** *Attack of the mutant tags*



<https://sysdig.com/blog/toctou-tag-mutability/>

<https://www.youtube.com/watch?v=j8K6EjOPhxs>




# Best practices - run

# Vulnerability management

- New vulnerabilities are **discovered daily**



- **Don't:** Use outdated images
  - Stick to the latest security patches
  - No need to always go with the latest version (breaking changes)



# Vulnerability management

- Runtime vulnerability assessment is key
  - **Re-evaluate** scanned images to detect new applying vulnerabilities
- **Update often + define a versioning strategy**
  - Stick to stable or long-term support versions
  - Rebuild periodically
    - To get the latest packages from the base distro
    - npm, go mod offer ways to specify version ranges (keep up with latest security updates)
  - Plan in advance.
    - Be ready to drop old versions
    - Migrate before base image reaches the end of its life and stops receiving updates

# Vulnerability management

- Fixing every single vulnerability might my **mission impossible**
- **Evaluate risk**: severity + exploitability + environment metrics
  - Is the environment critical (i.e. production)?
  - Does an exploit exist?
  - Is the exploit remote or local? Privileged or standard user? ...
  - Is the application exposed to the internet?
- Improve management and focus: use **alerting** and **ticketing** systems
- Upgrade or patch vs protect (i.e. firewall, configuration, ...)

# Runtime detection

- **Runtime threat detection** is critical. Falco can help.
- Create **specific** rules to detect exploitation of known vulnerabilities
- Generic rules to detect suspicious activity
  - i.e. unexpected processes, network connections, filesystem read/write, etc.
- Monitor & alert unexpected patterns
  - i.e. high CPU usage caused by cryptomining activities
- **Respond** to suspicious activity. i.e. stop or pause container + **forensics**



# Q & A